

GitSummary

I.Z.

September 29, 2018

```
knitr::opts_chunk$set(eval = FALSE, error = FALSE)
```

Record changes

1. Visualize changes:

In the shell, use

```
git status
git diff
```

to see an overview of changes and to show detailed differences. In RStudio click on Diff. The background colours tells you whether the text has been added (green) or removed (red). (If you're colourblind you can use the line numbers in the two columns at the far left as a guide: a number in the first column identifies the old version, a number in second column identifies the new version.) The grey lines of code above and below the changes give you additional context.

2. Commits:

The fundamental unit of work in Git is a commit. A commit takes a snapshot of your code at a specified point in time. Using a Git commit is like using anchors and other protection when climbing. If you're crossing a dangerous rock face you want to make sure you've used protection to catch you if you fall. Commits play a similar role: if you make a mistake, you can't fall past the previous commit. Coding without commits is like free-climbing: you can travel much faster in the short-term, but in the long-term the chances of catastrophic failure are high! Like rock climbing protection, you want to be judicious in your use of commits. Committing too frequently will slow your progress; use more commits when you're in uncertain or dangerous territory. Commits are also helpful to others, because they show your journey, not just the destination.

2.A. Creating commits:

- There are five key components to every commit:
 - A *unique identifier*, called a SHA (short for secure hash algorithm).
 - A *changeset* that describes which files were added, modified and deleted.
 - A human-readable *commit message*.
 - A *parent*, the commit that came before this one. (There are two exceptions to this rule: the initial commit doesn't have a parent, and merges, which you'll learn about later, have two parents.)
 - An *author*.

You create a commit in two stages: 1. You **stage** files, telling Git which changes should be included in the next commit. 1. You **commit** the staged files, describing the changes with a message.

In RStudio, staging and committing are done in the same place, the commit window, which you can open by pressing Ctrl + Alt + m.

The commit window is made up of three panes:

1. Top-left pane: current status as the Git pane in the main RStudio window
2. Bottom pane: shows the diff of the currently selected file (exactly the same window you see when clicking Diff)
3. Top-right pane: commit message

To create a new commit:

1. Save your changes.
2. Open the commit window by pressing Ctrl + Alt + m.
3. Select files: to stage (select) a single file for inclusion, tick its check box, to stage all files press Ctrl/Cmd + A, then click stage. As you stage each file, you'll notice that its status changes. The icon will change columns from right (unstaged status) to left (staged status), and you might see one of two new icons:
 - Added: after staging an untracked file, Git now knows that you want to add it to the repo.
 - Renamed: If you rename a file, Git initially sees it as a deletion and addition. Once you stage both changes, Git recognises that it's a rename. Sometimes you'll see a status in both columns. This means that you have both staged and unstaged changes in the same file. This happens when you've made some changes, staged them, and then made some more. Clicking the staged checkbox will stage your new changes, clicking it again will unstage both sets of changes.
4. Stage files, as above.
5. Write a commit message (top-right panel) which describes the changes that you've made. The first line of a commit message is called the subject line and should be brief (50 characters or less). For complicated commits, you can follow it with a blank line and then a paragraph or bulleted list providing more detail. Write messages in imperative, like you're telling someone what to do: "fix this bug", not "fixed this bug" or "this bug was fixed".
6. Click Commit.

Staging files is a little more complicated in the shell. You use `git add` to stage new and modified files, and `git rm` to stage deleted files. To create the commit, use `git commit -m <message>`.

2.B. Commit best practices:

Ideally, each commit should be minimal but complete:

- **Minimal:** A commit should only contain changes related to a single problem. This will make it easier to understand the commit at a glance, and to describe it with a simple message. If you should discover a new problem, you should do a separate commit.
- **Complete:** A commit should solve the problem that it claims to solve. If you think you've fixed a bug, the commit should contain a unit test that confirms you're right.

Each commit message should:

- **Be concise, yet evocative:** At a glance, you should be able to see what a commit does. Yet there should be enough detail so you can remember (and understand) what was done.
- **Describe the why, not the what:** Since you can always retrieve the diff associated with a commit, the message doesn't need to say exactly what changed. Instead it should provide a high-level summary that focuses on the reasons for the change.

If you do this:

- It'll be easier to work with others. For example, if two people have changed the same file in the same place, it'll be easier to resolve conflicts if the commits are small and it's clear why each change was made.
- Project newcomers can more easily understand the history by reading the commit logs.
- You can load and run your package at any point along its development history. This can be tremendously useful with tools like `bisectr`, which allow you to use binary search to quickly find the commit that introduced a bug.

- If you can figure out exactly when a bug was introduced, you can easily understand what you were doing (and why!).

You might think that because no one else will ever look at your repo, that writing good commit messages is not worth the effort. But keep in mind that you have one very important collaborator: future-you! If you spend a little time now polishing your commit messages, future-you will thank you if and when they need to do a post-mortem on a bug.

Remember that these directives are aspirational. You shouldn't let them get in your way. If you look at the commit history of my repositories, you'll notice a lot of them aren't that good, especially when I start to get frustrated that I still haven't managed to fix a bug. Strive to follow these guidelines, and remember it's better to have multiple bad commits than to have one perfect commit.

2.C. Ignoring files

Often, there are files that you don't want to include in the repository. They might be transient (like LaTeX or C build artefacts), very large, or generated on demand. Rather than carefully not staging them each time, you should instead add them to `.gitignore`. This will prevent them from accidentally being added. The easiest way to do this is to right-click on the file in the Git pane and select Ignore. If you want to ignore multiple files, you can use a wildcard "glob" like `*.png`.

Some developers never commit derived files, files that can be generated automatically. For an R package this would mean ignoring the files in the `NAMESPACE` and `man/` directories because they're generated from comments. From a practical perspective, it's better to commit these files: R packages have no way to generate `.Rd` files on installation so ignoring derived files means that users who install your package from GitHub will have no documentation.

blabla this is clearly a mistake